

# Introduction to RNA-seq for Differential Expression Analysis – 1

by Jiajia Li  
Biological Data Science Institute  
5 May 2025



Australian  
National  
University





## About the trainer – Jiajia Li

Bioinformatics IT Support Officer at RSB and BDSI

- Master of Data Science (AI and Computational Modelling)
  - University of Canberra, Australia
- Bachelor of Engineering (Bioengineering)
  - Henan University, China

Past Workshops:

- Introduction to Linux and Variant Calling
- Introduction to Python
- Data Visualisation with Python
- Machine Learning with Python
- Introduction to Git and GitHub
- Introduction to NCI Gadi Supercomputer
- Introduction to Snakemake Workflow Language
- Data Management and Reproducible Research

Any workshop you are interested in so I can run it again in Semester 2?!!!



# Workshop Goals

Due to the time limit and my capability, this workshop won't go deep into the interpretation of results. But we will still do a little bit...

We will focus on understanding the **Differential Gene Expression** workflow and get the code running!! (the most exciting part...)

## Bioinformatics and Biostatistics Drop-in Session

- Run **weekly** on **Tuesday** from **10-11am**, alternating locations between Robertson Building and JCSMR (shoot me an email if you don't know where it's on this week).
- If you have further questions about your own experiment, please welcome to join us at the drop-in session. We have an RNA-seq expert **Dr Zhi-Ping Feng** who can also provide help and advice.



# Reference

<https://rnabio.org/>



This is a great resource from the Griffith Lab at the Washington University. I developed this workshop based on their material. Their course has a lot of useful information from theory to analysis and interpretation. My workshop is a simplified version of it.

It is a good resource to look at if you want a more comprehensive understanding of RNA-sequencing techniques.

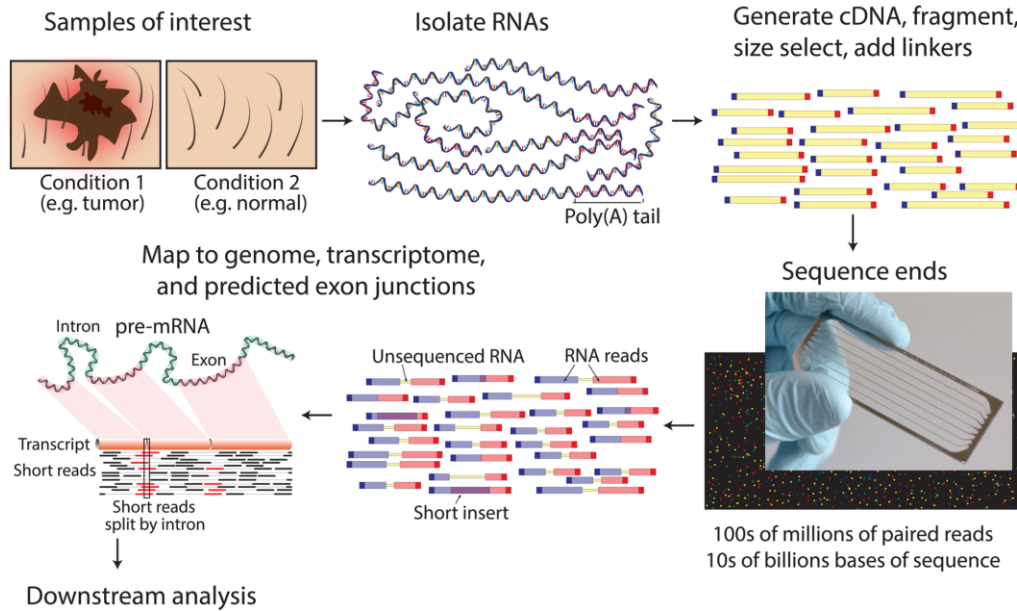
<https://www.nature.com/articles/s41576-019-0150-2> This is a great paper for introductory level of RNA sequencing.

# Learning Objectives of Today

- RNA-seq Analysis steps
- Different software
- Learn about our data
- Download data
- Learn about FASTQ/FASTA/GTF file format
- Learn about our reference genome
- Learn about Alignment
- Learn about Indexing
- Learn about software HISAT2
- Create an index for our reference genome
- Run FastQC and MultiQC on our data



# RNA-sequencing General Steps



Step1. Experimental Design (very important)

Step2. Wet Lab

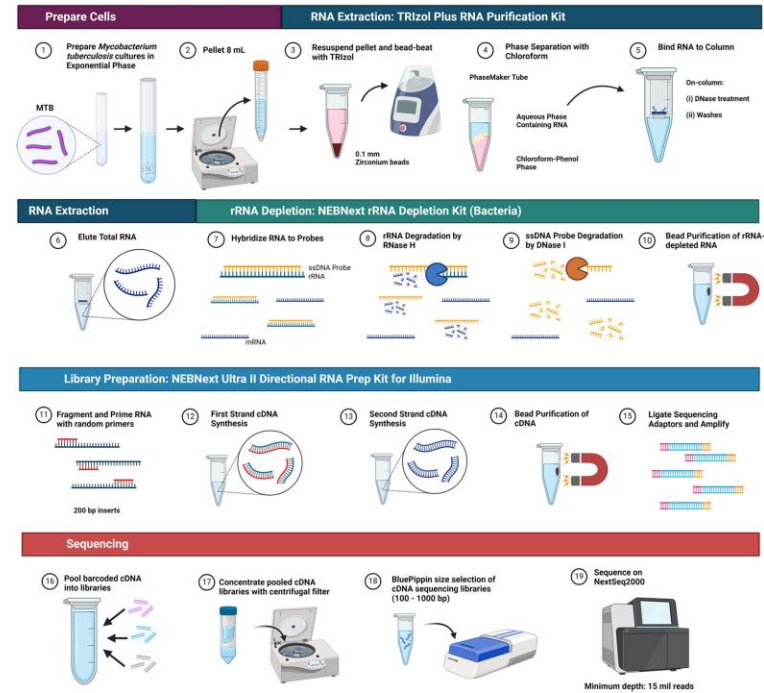
Step3. Dry Lab

# Wet Lab Steps

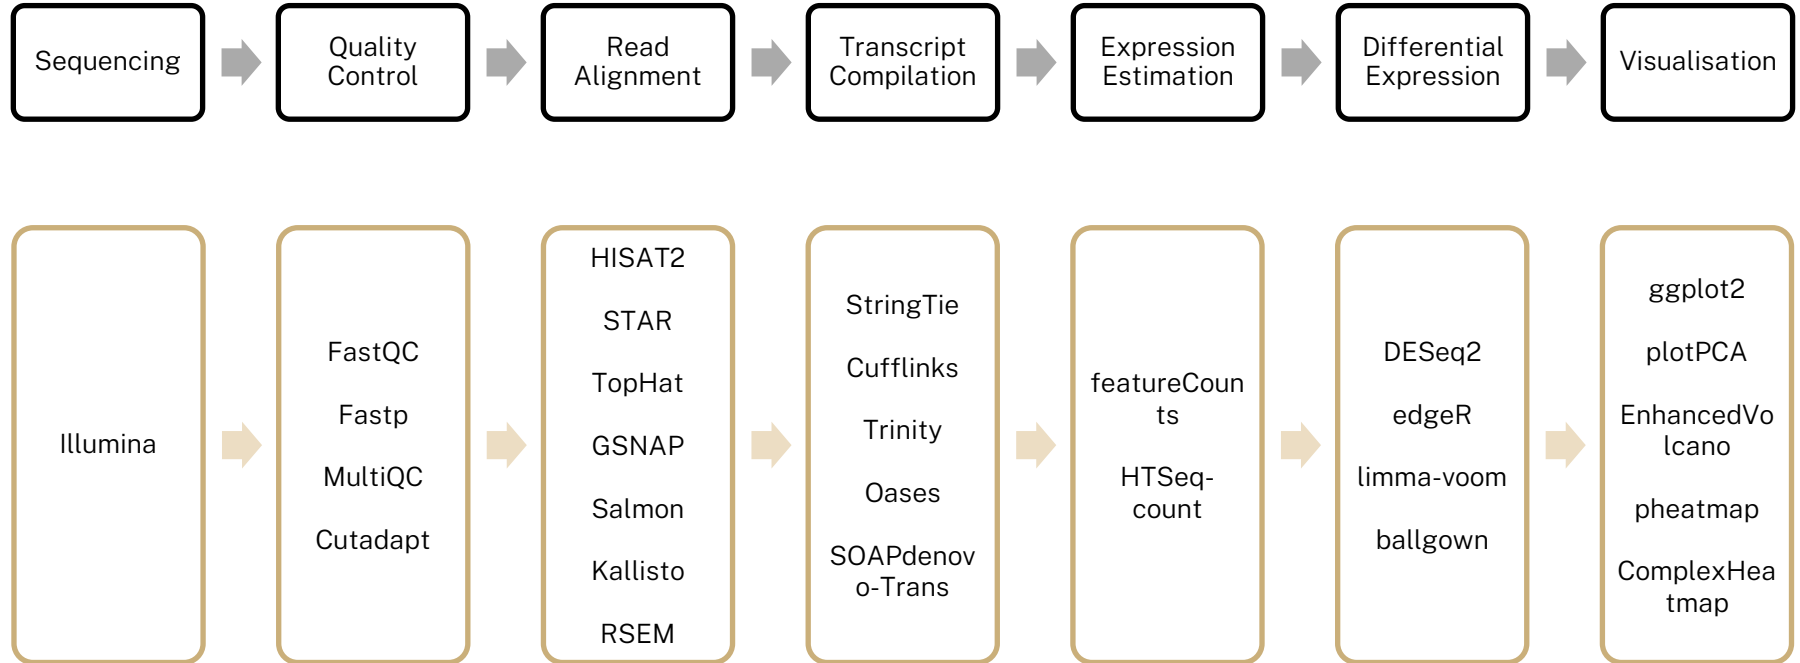
This might be done by the sequencing provider (e.g., BRF).

- 1) RNA extraction
- 2) RNA quality control
- 3) mRNA enrichment or rRNA depletion
- 4) RNA fragmentation
- 5) cDNA synthesis
- 6) Library preparation
- 7) Library quality control
- 8) Sequencing

Unfortunately, I won't be able to guide you in these steps...

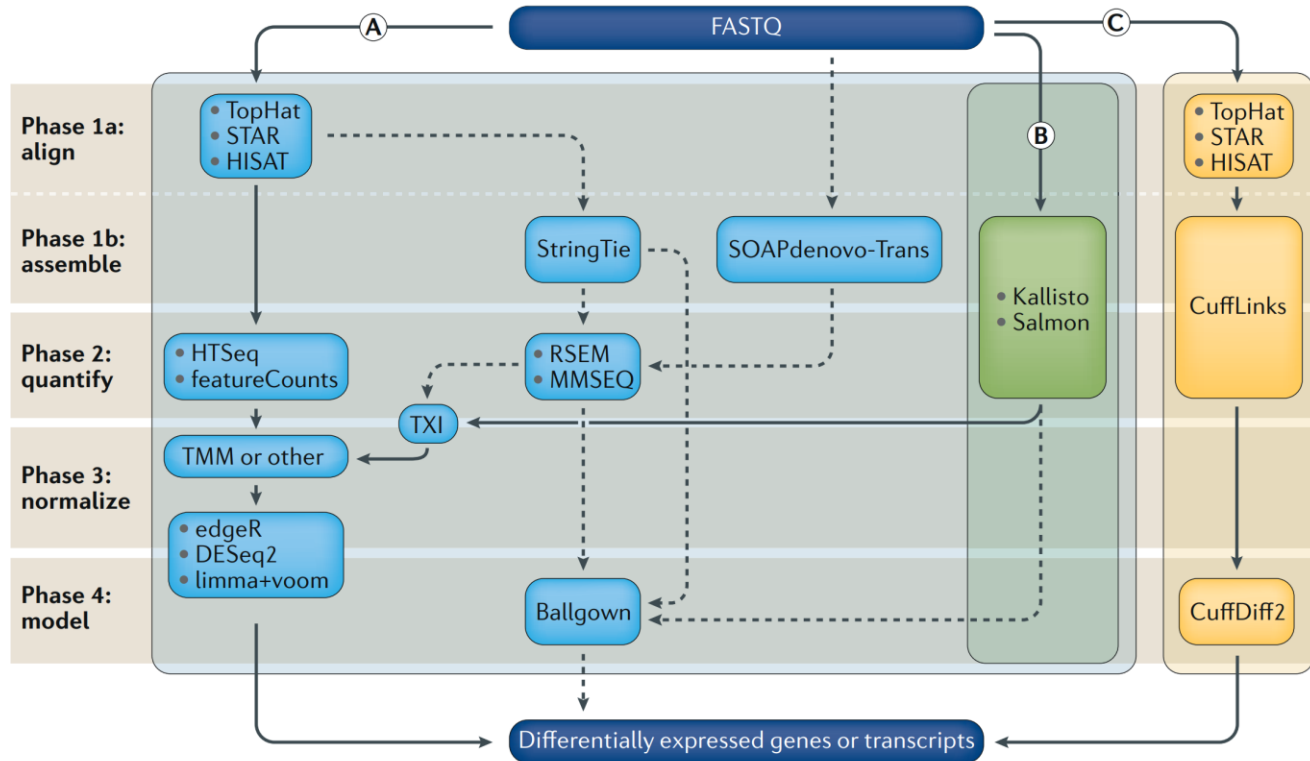


# Analysis Steps – Differential Expression





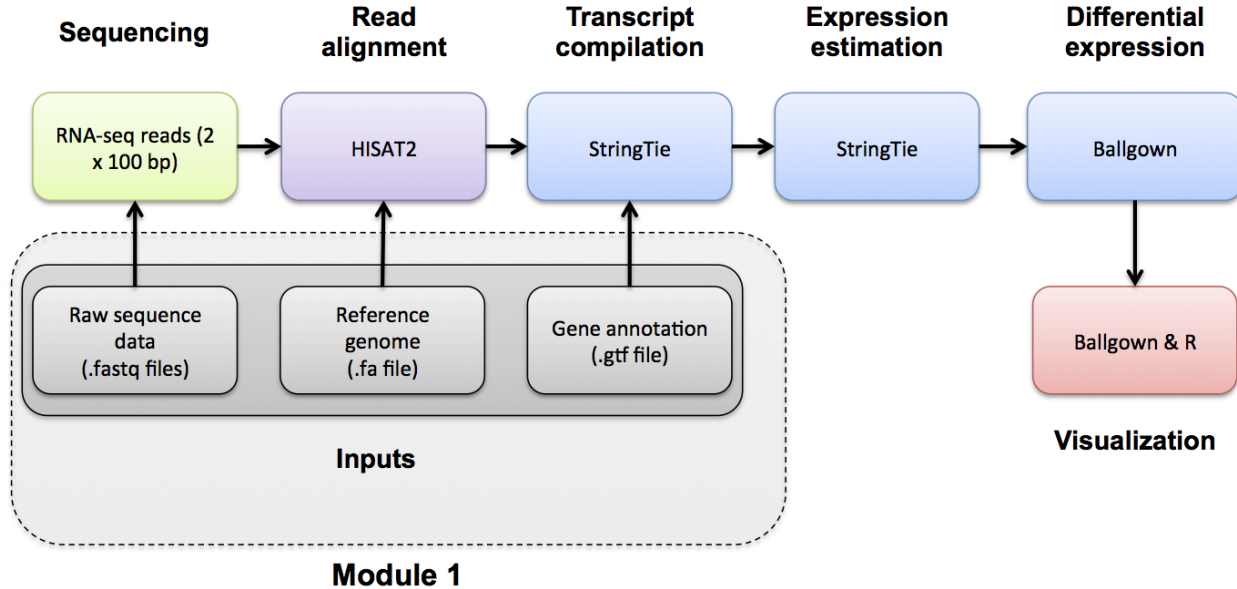
# Analysis Steps – Differential Expression



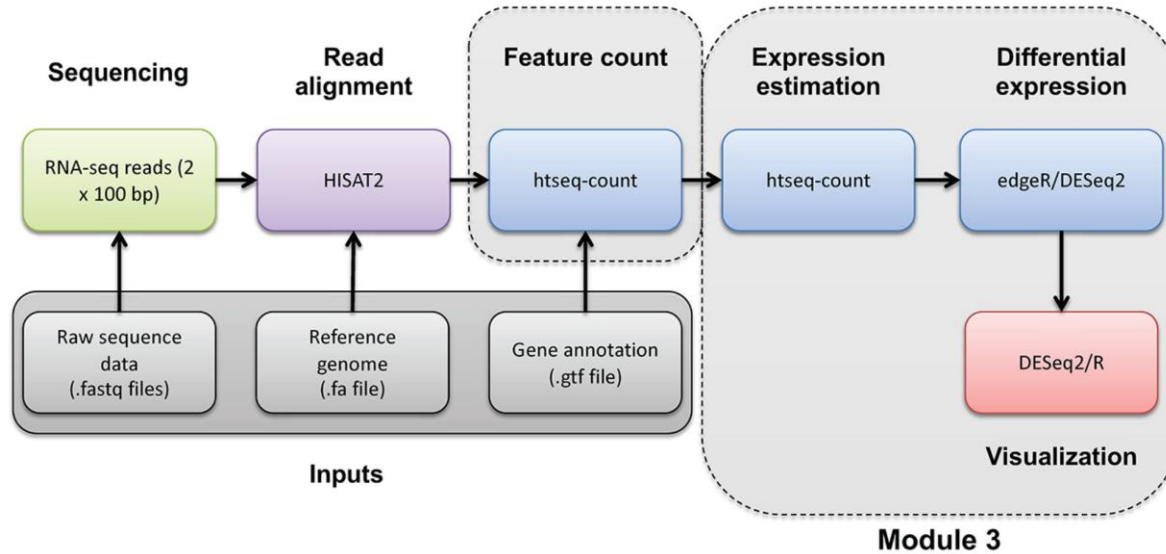
Possible combinations of pipelines.



# Pipeline 1 – HISAT2/StringTie/Ballgown



# Pipeline 2 – HISAT2/htseq-count/edgeR/DESeq2



# The data

The test data consists of two commercially available RNA samples:

- Universal Human Reference (UHR)
- Human Brain Reference (HBR)

The UHR is total RNA isolated from a diverse set of **10 cancer cell lines** (breast, liver, cervix, testis, brain, skin, fatty tissue, histocyte, macrophage, T cell, B cell).


The HBR is total RNA isolated from the **brain of 23 Caucasians**, male and female, of varying age but mostly 60-80 years old.

In addition, a spike-in control was used. An aliquot of the ERCC ExFold RNA Spike-In Control Mixes was added to each sample.



# ERCC ExFold RNA Spike-In Control Mixes

The spike-in consists of **92 transcripts** that are presented in known concentrations across a wide abundance range (from very few copies to many copies).



This range allows us to test the degree to which the RNA-seq assay (including all laboratory and analysis steps) accurately reflects the relative abundance of transcript species within a sample.

Because the concentration of transcripts in the spike-in control is known, if your analysis result couldn't correctly reflect the concentration of the spike-in control, it means your experiment is not reliable.




# ERCC ExFold RNA Spike-In Control Mixes

There are 2 “mixes” of these transcripts to allow an assessment of differential expression output between samples if you put one mix in each of your two comparisons.

In our case, Mix1 was added to the UHR sample, and Mix2 was added to the HBR sample.

For example, Mix1 has more transcript A than Mix2. If the result shows the UHR sample has less transcript A than the HBR sample, it means your experiment is not reliable.

We also have 3 complete experiment replicates for each sample. This allows us to assess the technical variability of our overall process of producing RNA-seq data in the lab.



# Library Preparation

This might be done by the service provider but it's great if you understand the theory.

All libraries were prepared using low-throughput **TruSeq Stranded Total RNA Sample Prep Kit** libraries with Ribo-Zero Gold to **remove both cytoplasmic and mitochondrial rRNA**.

- Why low-throughput?
  - because we only have 6 samples
- TruSeq Stranded Total RNA Sample Prep Kit
  - Except total RNA, you can also use mRNA prep kit, then your library with only have mRNAs
  - Total RNA will contain mRNA, tRNA, rRNA, miRNA, lncRNA, etc.
- Why remove rRNA?
  - In cells, rRNA accounts for about 80%-90% of the total RNA
  - If we don't remove rRNA, a large portion of our sequencing reads will be wasted on sequencing rRNA sequences



# Library Preparation

Triplicate, indexed libraries were made starting with 100ng Agilent/Stratagene Universal Human Reference total RNA and 100ng Ambion Human Brain Reference total RNA.

- Indexed – means that each library (or each sample) has a unique “barcode” attached to it.
- This allows you to pool multiple libraries together in a single sequencing run and then separate or “demultiplex” the reads back into their respective samples based on their indexing during analysis.

The UHR replicates received 2ul of 1:1000 ERCC Mix 1.

The HBR replicates received 2ul of 1:1000 ERCC Mix 2.

The libraries were quantified with KAPA Library Quantification qPCR and adjusted to the appropriate concentration for sequencing.

- Too concentrated or too diluted could negatively affect the sequencing result.

The triplicate, indexed libraries were then pooled prior to sequencing. Each pool of three replicate libraries were sequenced across 2 lanes of a HiSeq 2000 using paired-end sequence chemistry with **100bp** read lengths.





# Samples

To summarise we have:

- UHR + ERCC Spike-In Mix1, Replicate 1
- UHR + ERCC Spike-In Mix1, Replicate 2
- UHR + ERCC Spike-In Mix1, Replicate 3
  
- HBR + ERCC Spike-In Mix2, Replicate 1
- HBR + ERCC Spike-In Mix2, Replicate 2
- HBR + ERCC Spike-In Mix2, Replicate 3

Each data set has a corresponding pair of FASTQ file (read1 and read2 of paired-end reads).



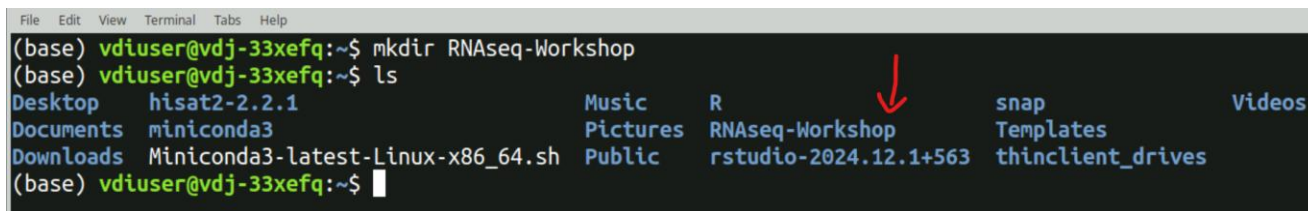
# Download Data

Now, let's go to <https://desktop.rc.nectar.org.au/> and open our ARDC Ubuntu machine.

Open Terminal and create a new directory “RNAseq-Workshop”.

To create a new folder in command line, we use command `mkdir RNAseq-Workshop`.

```
File Edit View Terminal Tabs Help
(base) vdiuser@vdj-33xefq:~$ mkdir RNAseq-Workshop
(base) vdiuser@vdj-33xefq:~$ ls
Desktop  hisat2-2.2.1      Music      R      snap      Videos
Documents miniconda3        Pictures   RNAseq-Workshop  Templates
Downloads Miniconda3-latest-Linux-x86_64.sh Public      rstudio-2024.12.1+563  thinclient_drives
(base) vdiuser@vdj-33xefq:~$
```



Go into this folder by running `cd RNAseq-Workshop`.

```
(base) vdiuser@vdj-33xefq:~$ cd RNAseq-Workshop/
(base) vdiuser@vdj-33xefq:~/RNAseq-Workshop$
```



# Download Data

Then, we will create another folder “data” inside the “RNAseq-Workshop” folder.

```
`mkdir data`  
(base) vdiuser@vdj-33xefq:~/RNAseq-Workshop$ mkdir data  
(base) vdiuser@vdj-33xefq:~/RNAseq-Workshop$ ls  
data
```

Go into this “data” folder, `cd data`.`

```
(base) vdiuser@vdj-33xefq:~/RNAseq-Workshop$ cd data  
(base) vdiuser@vdj-33xefq:~/RNAseq-Workshop/data$ ls  
(base) vdiuser@vdj-33xefq:~/RNAseq-Workshop/data$
```

Run `wget http://genomedata.org/rnaseq-tutorial/HBR_UHR_ERCC_ds_5pc.tar`` to download the .tar file, .tar is a type of compressed file.

```
(base) vdiuser@vdj-33xefq:~/RNAseq-Workshop/data$ wget http://genomedata.org/rnaseq-tutorial/HBR_UHR_ERCC_ds_5pc.  
tar  
--2025-04-27 12:20:06-- http://genomedata.org/rnaseq-tutorial/HBR_UHR_ERCC_ds_5pc.tar  
Resolving genomedata.org (genomedata.org)... 54.71.55.4  
Connecting to genomedata.org (genomedata.org)|54.71.55.4|:80... connected.  
HTTP request sent, awaiting response... 301 Moved Permanently  
Location: https://genomedata.org/rnaseq-tutorial/HBR_UHR_ERCC_ds_5pc.tar [following]  
--2025-04-27 12:20:06-- https://genomedata.org/rnaseq-tutorial/HBR_UHR_ERCC_ds_5pc.tar  
Connecting to genomedata.org (genomedata.org)|54.71.55.4|:443... connected.  
HTTP request sent, awaiting response... 200 OK  
Length: 116602880 (111M) [application/x-tar]  
Saving to: 'HBR_UHR_ERCC_ds_5pc.tar'  
  
HBR_UHR_ERCC_ds_5pc.tar 63%[=====] 71.07M 18.9MB/s eta 3s
```



# Download Data

Decompress the data by running  
``tar -xvf HBR_UHR_ERCC_ds_5pc.tar``

You should have total 12 files.

Check the size of your file using  
``ls -lh``

They seem to be small and around 6-13MB in size, that is because these reads have been pre-filtered to those only mapped to chromosome 22.

```
(base) vdiuser@vdj-33xefq:~/RNAseq-Workshop/data$ tar -xvf HBR_UHR_ERCC_ds_5pc.tar
HBR_Rep1_ERCC-Mix2_Build37-ErccTranscripts-chr22.read1.fastq.gz
HBR_Rep1_ERCC-Mix2_Build37-ErccTranscripts-chr22.read2.fastq.gz
HBR_Rep2_ERCC-Mix2_Build37-ErccTranscripts-chr22.read1.fastq.gz
HBR_Rep2_ERCC-Mix2_Build37-ErccTranscripts-chr22.read2.fastq.gz
HBR_Rep3_ERCC-Mix2_Build37-ErccTranscripts-chr22.read1.fastq.gz
HBR_Rep3_ERCC-Mix2_Build37-ErccTranscripts-chr22.read2.fastq.gz
UHR_Rep1_ERCC-Mix1_Build37-ErccTranscripts-chr22.read1.fastq.gz
UHR_Rep1_ERCC-Mix1_Build37-ErccTranscripts-chr22.read2.fastq.gz
UHR_Rep2_ERCC-Mix1_Build37-ErccTranscripts-chr22.read1.fastq.gz
UHR_Rep2_ERCC-Mix1_Build37-ErccTranscripts-chr22.read2.fastq.gz
UHR_Rep3_ERCC-Mix1_Build37-ErccTranscripts-chr22.read1.fastq.gz
UHR_Rep3_ERCC-Mix1_Build37-ErccTranscripts-chr22.read2.fastq.gz
```

```
(base) vdiuser@vdj-33xefq:~/RNAseq-Workshop/data$ ls -lh
total 223M
-rw-rw-r-- 1 vdiuser vdiuser 6.4M Nov 7 2014 HBR_Rep1_ERCC-Mix2_Build37-ErccTranscripts-chr22.read1.fastq.gz
-rw-rw-r-- 1 vdiuser vdiuser 6.7M Nov 7 2014 HBR_Rep1_ERCC-Mix2_Build37-ErccTranscripts-chr22.read2.fastq.gz
-rw-rw-r-- 1 vdiuser vdiuser 7.7M Nov 7 2014 HBR_Rep2_ERCC-Mix2_Build37-ErccTranscripts-chr22.read1.fastq.gz
-rw-rw-r-- 1 vdiuser vdiuser 8.1M Nov 7 2014 HBR_Rep2_ERCC-Mix2_Build37-ErccTranscripts-chr22.read2.fastq.gz
-rw-rw-r-- 1 vdiuser vdiuser 6.9M Nov 7 2014 HBR_Rep3_ERCC-Mix2_Build37-ErccTranscripts-chr22.read1.fastq.gz
-rw-rw-r-- 1 vdiuser vdiuser 7.3M Nov 7 2014 HBR_Rep3_ERCC-Mix2_Build37-ErccTranscripts-chr22.read2.fastq.gz
-rw-rw-r-- 1 vdiuser vdiuser 112M Oct 24 2018 HBR_UHR_ERCC_ds_5pc.tar
-rw-rw-r-- 1 vdiuser vdiuser 13M Nov 7 2014 UHR_Rep1_ERCC-Mix1_Build37-ErccTranscripts-chr22.read1.fastq.gz
-rw-rw-r-- 1 vdiuser vdiuser 14M Nov 7 2014 UHR_Rep1_ERCC-Mix1_Build37-ErccTranscripts-chr22.read2.fastq.gz
-rw-rw-r-- 1 vdiuser vdiuser 9.7M Nov 7 2014 UHR_Rep2_ERCC-Mix1_Build37-ErccTranscripts-chr22.read1.fastq.gz
-rw-rw-r-- 1 vdiuser vdiuser 11M Nov 7 2014 UHR_Rep2_ERCC-Mix1_Build37-ErccTranscripts-chr22.read2.fastq.gz
-rw-rw-r-- 1 vdiuser vdiuser 11M Nov 7 2014 UHR_Rep3_ERCC-Mix1_Build37-ErccTranscripts-chr22.read1.fastq.gz
-rw-rw-r-- 1 vdiuser vdiuser 12M Nov 7 2014 UHR_Rep3_ERCC-Mix1_Build37-ErccTranscripts-chr22.read2.fastq.gz
(base) vdiuser@vdj-33xefq:~/RNAseq-Workshop/data$
```



# Download Data

After we unzip the tar file, we don't need it anymore and can delete it by running  
``rm HBR_UHR_ERCC_ds_5pc.tar``

```
(base) vdiuser@vdj-33xefq:~/RNAseq-Workshop/data$ rm HBR_UHR_ERCC_ds_5pc.tar
(base) vdiuser@vdj-33xefq:~/RNAseq-Workshop/data$ ls
HBR_Rep1_ERCC-Mix2_Build37-ErccTranscripts-chr22.read1.fastq.gz
HBR_Rep1_ERCC-Mix2_Build37-ErccTranscripts-chr22.read2.fastq.gz
HBR_Rep2_ERCC-Mix2_Build37-ErccTranscripts-chr22.read1.fastq.gz
HBR_Rep2_ERCC-Mix2_Build37-ErccTranscripts-chr22.read2.fastq.gz
HBR_Rep3_ERCC-Mix2_Build37-ErccTranscripts-chr22.read1.fastq.gz
HBR_Rep3_ERCC-Mix2_Build37-ErccTranscripts-chr22.read2.fastq.gz
UHR_Rep1_ERCC-Mix1_Build37-ErccTranscripts-chr22.read1.fastq.gz
UHR_Rep1_ERCC-Mix1_Build37-ErccTranscripts-chr22.read2.fastq.gz
UHR_Rep2_ERCC-Mix1_Build37-ErccTranscripts-chr22.read1.fastq.gz
UHR_Rep2_ERCC-Mix1_Build37-ErccTranscripts-chr22.read2.fastq.gz
UHR_Rep3_ERCC-Mix1_Build37-ErccTranscripts-chr22.read1.fastq.gz
UHR_Rep3_ERCC-Mix1_Build37-ErccTranscripts-chr22.read2.fastq.gz
(base) vdiuser@vdj-33xefq:~/RNAseq-Workshop/data$
```



# FASTQ File

The FASTQ file is widely used to store both **sequence data** and its **quality scores**. A FASTQ file consists of **multiple entries**, with each entry representing a **single sequence read**.

Let's use our first file as an example to show you what a FASTQ file looks like.

Run `less HBR_Rep1_ERCC-Mix2_Build37-ErccTranscripts-chr22.read1.fastq.gz`

[illegible]

# FASTQ File

The command ``less`` is a file pager that allows you to **view the contents of a file** in a terminal, and it works with `.gz` files which is another format of compressed files. You can use the “↑” and “↓” arrow keys to navigate, or the “space” bar to go to next page.

We can also run `zcat HBR_Rep1_ERCC-Mix2_Build37-ErccTranscripts-chr22.read1.fastq.gz | head -n 4` to only print out the first 4 lines of the file on the terminal.

``cat`` is a command can print out all the content of a file to the screen, and ``zcat`` can print out the content of compressed files.

[illegible]



# FASTQ File

[illegible]

### Line 1 – Sequence Identifier

- Always start with an “@”.
- Metadata about the read, such as read name, sample info, or a unique identifier.

## Line 2 – Sequence Line

- The actual sequence, A, T, G, C

### Line 3 – Plus Sign “+”

- This line serves as a separator.

### Line 4 – Quality Scores

- A string of **ASCII characters** that represent the quality of each base in the sequence.
- Typically encoded using the **Phred quality score system**.





# FASTQ File

Count how many reads in your FASTQ file. Since we know the first line starts with @ and it represents a sequence read. If we count the number of lines that start with @ we would know how many reads in the file.

To do so, we can run

```
`zcat HBR_Rep1_ERCC-Mix2_Build37-ErccTranscripts-chr22.read1.fastq.gz | grep -c '^@`
```

- grep – a command search for patterns
- “^@” – search for lines start with @
- -c – count the number of lines that match the pattern

```
(base) vdiuser@vdj-33xefq:~/RNAseq-Workshop/data$ zcat HBR_Rep1_ERCC-Mix2_Build37-ErccTranscripts-chr22.read1.fastq.gz | grep -c '^@'  
144098  
(base) vdiuser@vdj-33xefq:~/RNAseq-Workshop/data$
```



# Reference Genome

In this analysis we will use the **GRCh38** version of the human genome from Ensemble ([https://ftp.ensembl.org/pub/release-86/fasta/homo\\_sapiens/dna/](https://ftp.ensembl.org/pub/release-86/fasta/homo_sapiens/dna/)). We are going to perform the analysis using only a single chromosome (**chr22**) and the ERCC spike-in to make it run faster.

Let's go back to our "RNAseq-Workshop" folder and create a new folder called "reference".

Use `cd ..` to move one level up.

```
(base) vdiuser@vdj-33xefq:~/RNAseq-Workshop/data$ cd ..  
(base) vdiuser@vdj-33xefq:~/RNAseq-Workshop$
```

Then run `mkdir reference`.

```
(base) vdiuser@vdj-33xefq:~/RNAseq-Workshop$ mkdir reference  
(base) vdiuser@vdj-33xefq:~/RNAseq-Workshop$ ls  
data  reference  
(base) vdiuser@vdj-33xefq:~/RNAseq-Workshop$
```



# Reference Genome

Go into the newly created “reference” folder - `cd reference``

Run `wget http://genomedata.org/rnaseq-tutorial/fasta/GRCh38/chr22_with_ERCC92.fa`` to download the references.

This is a simplified reference, contains only chr22 and ERCC transcripts.

```
(base) vdiuser@vdj-33xefq:~/RNAseq-Workshop$ cd reference/
(base) vdiuser@vdj-33xefq:~/RNAseq-Workshop/reference$ wget http://genomedata.org/rnaseq-tutorial/fasta/GRCh38/chr22_with_ERCC92.fa
--2025-04-27 14:14:42-- http://genomedata.org/rnaseq-tutorial/fasta/GRCh38/chr22_with_ERCC92.fa
Resolving genomedata.org (genomedata.org)... 54.71.55.4
Connecting to genomedata.org (genomedata.org)|54.71.55.4|:80... connected.
HTTP request sent, awaiting response... 301 Moved Permanently
Location: https://genomedata.org/rnaseq-tutorial/fasta/GRCh38/chr22_with_ERCC92.fa [following]
--2025-04-27 14:14:43-- https://genomedata.org/rnaseq-tutorial/fasta/GRCh38/chr22_with_ERCC92.fa
Connecting to genomedata.org (genomedata.org)|54.71.55.4|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 51751056 (49M)
Saving to: 'chr22_with_ERCC92.fa'

chr22_with_ERCC92.fa      3%[>] 1.93M 1.81MB/s
```





# FASTA File Format

FASTA is a widely used file format for representing **nucleotide sequences** (DNA or RNA) or protein sequences.

A FASTA file can contain multiple entries, each entry has the following structure:

## Header Line:

- Starts with an “>”
- Then description or identifier about the sequence.

## Sequence Lines:

- A, T, G, C/A, U, G, C/amino acids
- The sequence can be broken into **multiple lines** to avoid very long lines, but the sequence itself should be continuous without spaces, gaps or any other characters.



# FASTA File Format

Check how many sequences in our reference genome.

```
`grep -c '^>' chr22_with_ERCC92.fa`
```

```
(base) vdiuser@vdj-33xefq:~/RNAseq-Workshop/reference$ grep -c '^>' chr22_with_ERCC92.fa  
93
```

93 sequences!! why?? usually for a chromosome it should be only one single sequence.

We can search all the lines start with ">" and print them out to see.

```
`grep '^>' chr22_withERCC92.fa`
```

```
(base) vdiuser@vdj-33xefq:~/RNAseq-Workshop/reference$ grep '^>' chr22_with_ERCC92.fa  
>22 dna_sm:chromosome chromosome:GRCh38:22:1:50818468:1 REF  
>ERCC-00002  
>ERCC-00003  
>ERCC-00004  
>ERCC-00009  
>ERCC-00012  
>ERCC-00013  
>ERCC-00014  
>ERCC-00016
```



# Reference Genome

If you download the same chromosome 22 for GRCh38 from other resources, such as UCSC, the name shows in the “>” line could be different.

Remember that the names of your reference sequences (chromosomes) must match those in your annotation GTF files.



# Known Gene/Transcript Annotations

In this tutorial we will use annotations from Ensembl ([https://ftp.ensembl.org/pub/release-86/gtf/homo\\_sapiens/](https://ftp.ensembl.org/pub/release-86/gtf/homo_sapiens/)) for chromosome 22 only.

For time reasons, these are prepared for you.

Make sure you are in the “reference” folder and run:

```
`wget http://genomedata.org/rnaseq-tutorial/annotations/GRCh38/chr22_with_ERCC92.gtf`
```

```
(base) vdiuser@vdj-33xefq:~/RNAseq-Workshop/reference$ ls -lh
total 79M
-rw-rw-r-- 1 vdiuser vdiuser 50M Oct 24 2018 chr22_with_ERCC92.fa
-rw-rw-r-- 1 vdiuser vdiuser 30M Oct 24 2018 chr22_with_ERCC92.gtf
```

Use ``less chr22_with_ERCC92.gtf`` to have a look of the file content.

```
22   ensembl gene   10736171       10736283       -       gene_id "EN
SG00000277248"; gene_version "1"; gene_name "U2"; gene_source "ensembl"; gene_biotype "snRNA";
22   ensembl transcript 10736171       10736283       -       gene
e_id "ENSG00000277248"; gene_version "1"; transcript_id "ENST00000615943"; transcript_vers
on "1"; gene_name "U2"; gene_source "ensembl"; gene_biotype "snRNA"; transcript_name "U2.14
-201"; transcript_source "ensembl"; transcript_biotype "snRNA"; tag "basic"; transcript_sup
port_level "NA";
22   ensembl exon   10736171       10736283       -       gene_id "EN
SG00000277248"; gene_version "1"; transcript_id "ENST00000615943"; transcript_version "1";
exon_number "1"; gene_name "U2"; gene_source "ensembl"; gene_biotype "snRNA"; transcript_na
me "U2.14-201"; transcript_source "ensembl"; transcript_biotype "snRNA"; exon_id "ENSE00003
736336"; exon_version "1"; tag "basic"; transcript_support_level "NA";
22   havana gene   10939388       10961338       -       gene_id "EN
SG00000283047"; gene_version "1"; gene_name "FRG1FP"; gene_source "havana"; gene_biotype "u
nprocessed_pseudogene"; havana_gene "OTTHUMC00000191577"; havana_gene_version "1";
chr22_with_ERCC92.gtf
```





# Known Gene/Transcript Annotations

It looks a bit messy because each line is too long, and it returns to the second line.

Use ``less -S chr22_with_ERCC92.gtf`` to view it in one line and in a left-right scrollable way.

```
22   ensembl gene      10736171      10736283      .      -      .      gene_id >
22   ensembl transcript 10736171      10736283      .      -      .      .      >
22   ensembl exon      10736171      10736283      .      -      .      gene_id >
22   havana gene      10939388      10961338      .      -      .      gene_id >
22   havana transcript 10939388      10961338      .      -      .      .      >
22   havana exon      10961283      10961338      .      -      .      gene_id >
22   havana exon      10959067      10959136      .      -      .      gene_id >
22   havana exon      10950049      10950174      .      -      .      gene_id >
22   havana exon      10949212      10949269      .      -      .      gene_id >
22   havana exon      10947304      10947418      .      -      .      gene_id >
22   havana exon      10944967      10945053      .      -      .      gene_id >
22   havana exon      10941691      10941780      .      -      .      gene_id >
22   havana exon      10940597      10940707      .      -      .      gene_id >
22   havana exon      10939388      10939423      .      -      .      gene_id >
22   ensembl gene      11065974      11067346      .      -      .      gene_id >
22   ensembl transcript 11065974      11067346      .      -      .      .      >
:|
```

But still not very good...

Try ``cat chr22_with_ERCC92.gtf | column -t | less -p exon -S``, it looks better.



# GTF File Format

GTF stands for **Gene Transfer Format** (sometimes called GFF version 2.5).

It's a text file used in genomics to describe **gene structures** – like where genes, exons, transcripts are located on a genome.

Each line in a GTF file describes one feature (such as an exon, a transcript, a CDS, etc.).

Fields are tab-separated – **9 fields per line**.

```
22      ensembl      transcript  10736171  10736283  .      -      .  gene_id>
22      ensembl      exon        10736171  10736283  .      -      .  gene_id>
22      havana       gene        10939388  10961338  .      -      .  gene_id>
22      havana       transcript  10939388  10961338  .      -      .  gene_id>
22      havana       exon        10961283  10961338  .      -      .  gene_id>
22      havana       exon        10959067  10959136  .      -      .  gene_id>
22      havana       exon        10950049  10950174  .      -      .  gene_id>
22      havana       exon        10949212  10949269  .      -      .  gene_id>
22      havana       exon        10947304  10947418  .      -      .  gene_id>
22      havana       exon        10944967  10945053  .      -      .  gene_id>
22      havana       exon        10941691  10941780  .      -      .  gene_id>
22      havana       exon        10940597  10940707  .      -      .  gene_id>
22      havana       exon        10939388  10939423  .      -      .  gene_id>
22      ensembl      gene        11065974  11067346  .      -      .  gene_id>
22      ensembl      transcript  11065974  11067346  .      -      .  gene_id>
22      ensembl      exon        11067335  11067346  .      -      .  gene_id>
:
```



# GTF File Format

Column	What it means	Example
1. Seqname	Chromosome or scaffold (chr1, chr22, MT, ...)	chr22
2. Source	Where the annotation comes from (Ensembl, HAVANA, ...)	ensembl
3. Feature	Types of feature (gene, transcript, exon, CDS, ...)	exon
4. Start	Start position	10736171
5. End	End position (inclusive)	10736283
6. Score	A score (like confidence), can be . if not used	.
7. Strand	+ or - (which DNA strand)	-
8. Frame	Reading frame (0, 1, 2), or . is not applicable	.
9. Attributes	Extra information in key-value pairs, separated by semicolon “;”	gene_id "ENSG00000277248"; gene_version "1"; transcript_id "ENST00000615943"; ...



# Indexing Reference Genome


Indexing a reference genome is to prepare it for alignment tools can **find matching regions quickly**.

Think of it like:

- Without an index – searching for a sequence in the whole genome would be like reading through an entire book page by page.
- With an index – you have a table of contents and page numbers; you can jump directly to the right place.

Because genomes are normally huge, and it is very slow search the whole thing every time.





HISAT2 stands for Hierarchical Indexing for Spliced Alignment of Transcripts (version 2). It is a very fast and efficient aligner for **mapping short sequencing reads (like Illumina) to a reference genome.**

It's especially good at:

- Handling **RNA-seq data** (reads that span **splice junctions**)
- Aligning to **large genomes** like human or mouse

It's kind of the next generation tools after older aligners like **TopHat** and **Bowtie**.



# Alignment

Alignment means – “Mapping sequencing reads back to a reference genome (or transcriptome) to find out **where they came from**”.

- In DNA-seq: reads are mapped **directly** to the genomic DNA. Straightforward and no introns to skip.
- In RNA-seq: reads are mapped to the genome, but they might **skip introns**, because RNA comes from **spliced transcripts**. Complex because reads may span exon-exon junctions.

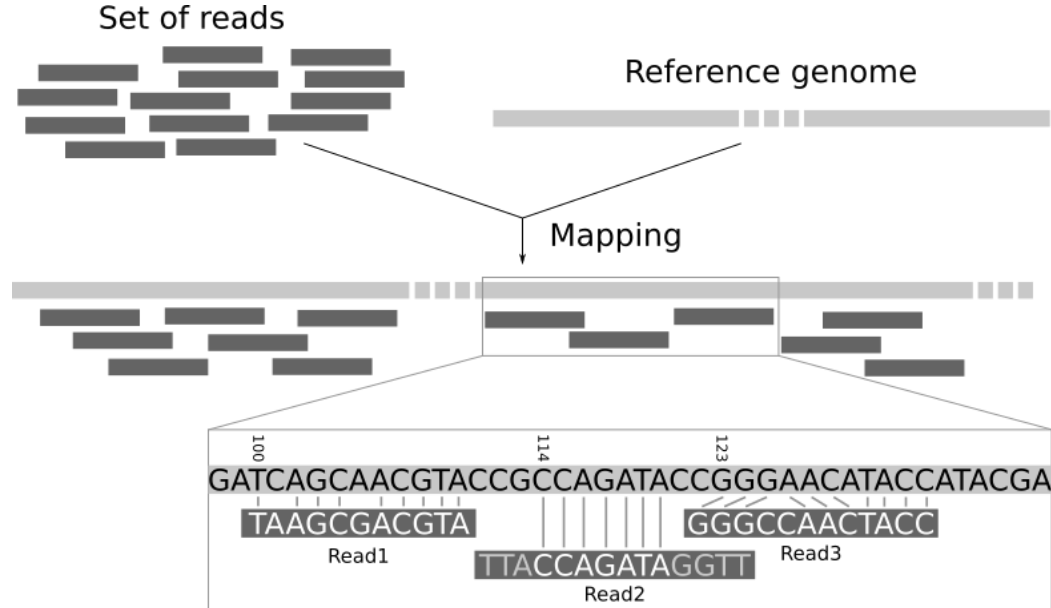
You could map RNA-seq to a transcriptome directly as well, if you have a known transcriptome. Tools like **Salmon** and **Kallisto** can do it.

You could also perform **de novo transcriptome assembly** which assemble transcripts directly from the RNA-seq reads. Tools such as Trinity, rnaSPAdes, etc. can do it.



# Alignment


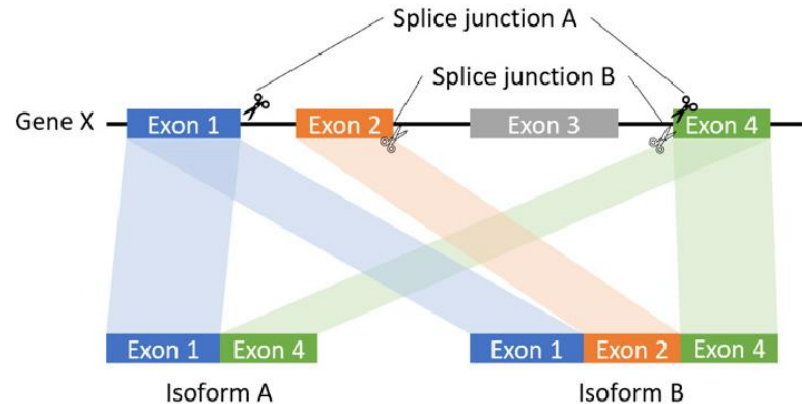
Illustration for DNA-seq mapping. For each read it normally fully matching to the reference genome.



# Alignment

For an RNA sequencing read, it could possibly be cut by introns and when it maps to genome, there will be a gap.

It's not necessarily that each of your read could be cut by introns, if the exons are longer than your read length, it can stay intact.





# Create a HISAT2 index

Now let's create an index for our reference, here we need to use a few python scripts from the HISAT2 software.

There are 3 steps:

- Extract splice sites
- Extract exons
- Build the index

First, in your ARDC ubuntu machine, open Terminal. And navigate to where your reference located.

``cd ~/RNAseq-Workshop/reference``

```
(base) vdiuser@vdj-33xefq:~$ cd ~/RNAseq-Workshop/reference/  
(base) vdiuser@vdj-33xefq:~/RNAseq-Workshop/reference$
```



# Create a HISAT2 index

Then run `~/hisat2-2.2.1/extract_splice_sites.py chr22_with_ERCC92.gtf > splicesites.tsv` to extract the splice sites from genome annotation file.

```
(base) vdiuser@vdj-33xefq:~/RNAseq-Workshop/reference$ ~/hisat2-2.2.1/extract_splice_sites.py chr22_with_ERCC92.gtf > splicesites.tsv
(base) vdiuser@vdj-33xefq:~/RNAseq-Workshop/reference$
```

Then run `~/hisat2-2.2.1/extract_exons.py chr22_with_ERCC92.gtf > exons.tsv` to extract exons.

```
(base) vdiuser@vdj-33xefq:~/RNAseq-Workshop/reference$ ~/hisat2-2.2.1/extract_exons.py chr22_with_ERCC92.gtf > exons.tsv
(base) vdiuser@vdj-33xefq:~/RNAseq-Workshop/reference$
```



# Create a HISAT2 index

Finally, we use the splice sites and exons information to build an index on the reference genome (FASTA file).

Run `~/hisat2-2.2.1/hisat2-build -p 4 --ss splicesites.tsv --exon exons.tsv chr22_with_ERCC92.fa chr22_with_ERCC92_index``

-p to specify how many threads to use.

--ss to specify splice sites file.

--exon to specify exons file.

Use ``ls`` to check your result files!!

```
(base) vdiuser@vdj-33xefq:~/RNAseq-Workshop/reference$ ls
chr22_with_ERCC92.fa      chr22_with_ERCC92_index.5.ht2
chr22_with_ERCC92.gtf     chr22_with_ERCC92_index.6.ht2
chr22_with_ERCC92_index.1.ht2 chr22_with_ERCC92_index.7.ht2
chr22_with_ERCC92_index.2.ht2 chr22_with_ERCC92_index.8.ht2
chr22_with_ERCC92_index.3.ht2 exons.tsv
chr22_with_ERCC92_index.4.ht2 splicesites.tsv
```

```
(base) vdiuser@vdj-33xefq:~/RNAseq-Workshop/reference$ ~/hisat2-2.2.1/hisat2-build -p 4 --ss splicesites.tsv --exon exons.tsv chr22_with_ERCC92.fa chr22_with_ERCC92_index
Settings:
  Output files: "chr22_with_ERCC92_index.*.ht2"
  Line rate: 7 (line is 128 bytes)
  Lines per side: 1 (side is 128 bytes)
  Offset rate: 4 (one in 16)
  FTable chars: 10
  Strings: unpacked
  Local offset rate: 3 (one in 8)
  Local fTable chars: 6
```

```
numSides: 189151
numLines: 189151
gbwtTotLen: 24211328
gbwtTotSz: 24211328
reverse: 0
linearFM: No
Total time for call to driver() for forward index: 00:01:02
(base) vdiuser@vdj-33xefq:~/RNAseq-Workshop/reference$
```



# Pre-alignment Quality Control

We can use **FastQC** to get a sense of your data quality before alignment. FastQC can generate a comprehensive report, showing several key metrics about the sequencing quality.

We have installed FastQC through Conda, so let's activate our Conda environment first.

```
`conda activate RNAseq_env`
```

```
(base) vdiuser@vdj-33xefq:~/RNAseq-Workshop/reference$ conda activate RNAseq_env  
(RNAseq_env) vdiuser@vdj-33xefq:~/RNAseq-Workshop/reference$
```

Then we go to the “data” folder where our sequencing files locates.

```
`cd ../data`
```

```
(RNAseq_env) vdiuser@vdj-33xefq:~/RNAseq-Workshop/reference$ cd ../data/  
(RNAseq_env) vdiuser@vdj-33xefq:~/RNAseq-Workshop/data$ ls  
HBR_Rep1_ERCC-Mix2_Build37-ErccTranscripts-chr22.read1.fastq.gz  
HBR_Rep1_ERCC-Mix2_Build37-ErccTranscripts-chr22.read2.fastq.gz  
HBR_Rep2_ERCC-Mix2_Build37-ErccTranscripts-chr22.read1.fastq.gz  
HBR_Rep2_ERCC-Mix2_Build37-ErccTranscripts-chr22.read2.fastq.gz  
HBR_Rep3_ERCC-Mix2_Build37-ErccTranscripts-chr22.read1.fastq.gz  
HBR_Rep3_ERCC-Mix2_Build37-ErccTranscripts-chr22.read2.fastq.gz  
UHR_Rep1_ERCC-Mix1_Build37-ErccTranscripts-chr22.read1.fastq.gz  
UHR_Rep1_ERCC-Mix1_Build37-ErccTranscripts-chr22.read2.fastq.gz  
UHR_Rep2_ERCC-Mix1_Build37-ErccTranscripts-chr22.read1.fastq.gz  
UHR_Rep2_ERCC-Mix1_Build37-ErccTranscripts-chr22.read2.fastq.gz  
UHR_Rep3_ERCC-Mix1_Build37-ErccTranscripts-chr22.read1.fastq.gz  
UHR_Rep3_ERCC-Mix1_Build37-ErccTranscripts-chr22.read2.fastq.gz
```



# Pre-alignment Quality Control

Then we can run `fastqc *.fastq.gz`

“\*” is a **wildcard** in Linux, it will match all the files that end with `.fastq.gz` under your current directory.

```
(RNAseq_env) vdiuser@vdj-33xefq:~/RNAseq-Workshop/data$ fastqc *.fastq.gz
application/gzip
application/gzip
application/gzip
Started analysis of HBR_Rep1_ERCC-Mix2_Build37-ErccTranscripts-chr22.read1.fastq.
gz
application/gzip
application/gzip
application/gzip
application/gzip
application/gzip
application/gzip
application/gzip
application/gzip
Approx 5% complete for HBR_Rep1_ERCC-Mix2_Build37-ErccTranscripts-chr22.read1.fastq.gz
Approx 10% complete for HBR_Rep1_ERCC-Mix2_Build37-ErccTranscripts-chr22.read1.fastq.gz
Approx 15% complete for HBR_Rep1_ERCC-Mix2_Build37-ErccTranscripts-chr22.read1.fastq.gz
```

Use `ls` to check the result.

```
(RNAseq_env) vdiuser@vdj-33xefq:~/RNAseq-Workshop/data$ ls
HBR_Rep1_ERCC-Mix2_Build37-ErccTranscripts-chr22.read1_fastqc.html
HBR_Rep1_ERCC-Mix2_Build37-ErccTranscripts-chr22.read1_fastqc.zip
HBR_Rep1_ERCC-Mix2_Build37-ErccTranscripts-chr22.read1.fastq.gz
HBR_Rep1_ERCC-Mix2_Build37-ErccTranscripts-chr22.read2_fastqc.html
HBR_Rep1_ERCC-Mix2_Build37-ErccTranscripts-chr22.read2_fastqc.zip
HBR_Rep1_ERCC-Mix2_Build37-ErccTranscripts-chr22.read2.fastq.gz
HBR_Rep2_ERCC-Mix2_Build37-ErccTranscripts-chr22.read1_fastqc.html
HBR_Rep2_ERCC-Mix2_Build37-ErccTranscripts-chr22.read1_fastqc.zip
HBR_Rep2_ERCC-Mix2_Build37-ErccTranscripts-chr22.read1.fastq.gz
```




# Pre-alignment Quality Control

For each file, there is a HTML, and a ZIP file generated. The HTML can be opened by a web browser and will contain a report for this FASTQ file.

On your Desktop, open the Home folder, and navigate to where these HTML reports located. It should be under “**RNAseq-Workshop/data**”, then you can double click to open.

Let's open the first one “HBR\_Rep1\_read1.HTML”

 **FastQC Report** Mon. 28 Apr. 2025  
HBR\_Rep1\_ERCC-Mix2\_Build37-ErccTranscripts-chr22.read1.fastq.gz

**Summary**

- ✓ [Basic Statistics](#)
- ✓ [Per base sequence quality](#)
- ✓ [Per tile sequence quality](#)
- ✓ [Per sequence quality scores](#)
- ✗ [Per base sequence content](#)
- ! [Per sequence GC content](#)
- ✓ [Per base N content](#)
- ✓ [Sequence Length Distribution](#)
- ! [Sequence Duplication Levels](#)
- ! [Overrepresented sequences](#)
- ✓ [Adapter Content](#)

**Basic Statistics**

Measure	Value
Filename	HBR_Rep1_ERCC-Mix2_Build37-ErccTranscripts-chr22.read1.fastq.gz
File type	Conventional base calls
Encoding	Sanger / Illumina 1.9
Total Sequences	118571
Total Bases	11.8 Mbp
Sequences flagged as poor quality	0
Sequence length	100
%GC	50

**Per base sequence quality**


Quality scores across all bases (Sanger / Illumina 1.9 encoding)



# Pre-alignment Quality Control

On the left pane, there are a few metrics. It can give you an overall idea of how your sample data looks like.












**Green** – Pass. **Red** – Fail. **Yellow** – Warning.


 **FastQC Report**

HBR\_Rep1\_ERCC-Mix2\_Build37-ErccTranscripts-chr22.read1.fastq.gz


Mon. 28 Apr. 2025

**Summary**

-  [Basic Statistics](#)
-  [Per base sequence quality](#)
-  [Per tile sequence quality](#)
-  [Per sequence quality scores](#)
-  [Per base sequence content](#)
-  [Per sequence GC content](#)
-  [Per base N content](#)
-  [Sequence Length Distribution](#)
-  [Sequence Duplication Levels](#)
-  [Overrepresented sequences](#)
-  [Adapter Content](#)

 **Basic Statistics**

Measure	Value
Filename	HBR_Rep1_ERCC-Mix2_Build37-ErccTranscripts-chr22.read1.fastq.gz
File type	Conventional base calls
Encoding	Sanger / Illumina 1.9
Total Sequences	118571
Total Bases	11.8 Mbp
Sequences flagged as poor quality	0
Sequence length	100
%GC	50

 **Per base sequence quality**

Quality scores across all bases (Sanger / Illumina 1.9 encoding)

Our sample looks alright;  
we have 1 failing and 3  
warnings.



# Pre-alignment Quality Control

Here is a link of what a good Illumina short-read sequence should look like.

[https://www.bioinformatics.babraham.ac.uk/projects/fastqc/good\\_sequence\\_short\\_fastqc.html](https://www.bioinformatics.babraham.ac.uk/projects/fastqc/good_sequence_short_fastqc.html)

We are not going to go through all the matrices, but we will look at those with issues.



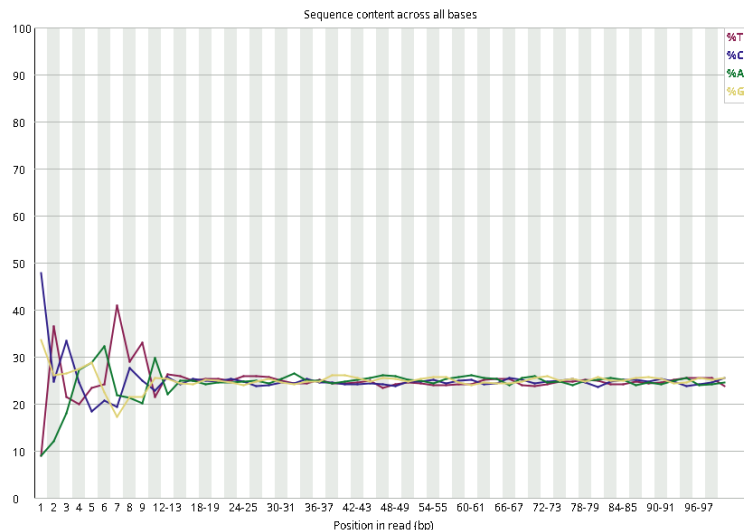


# Per base sequence content

For each position across all your reads, it plots the **percentage** of A, T, G, and C bases at that position.

For example, on position 1 (base 1), there are 10% A, 10% T, 35% G, and 50% C.

## Per base sequence content



In an ideal case, if your library is random, you expect the proportions of A, T, G, C, to be **pretty flat** and **similar** across positions.

They should all be around 25% each.

This sample seems to have **small wobbles** at the beginning but becomes stable after **13** bases which is good.

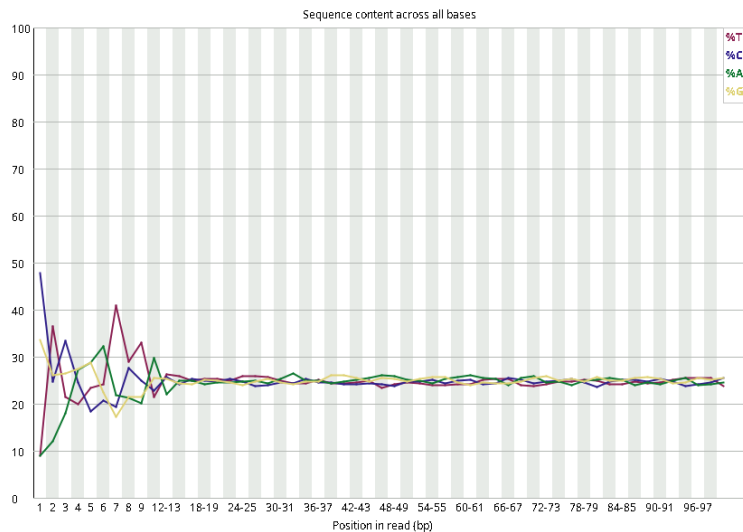


# Per base sequence content

Should we trim it?

Some biases are normal at the very start (1-10 bases) due to primers or **technical effects**, especially in RNA-seq or amplicon sequencing.

## ✖ Per base sequence content



If the wobbles are small and settle down after, you probably don't have to trim it.


Many people just leave it because modern aligners (like HISAT2) **handle small biases pretty well**.

But if you want to be super clean or your analysis is sensitive (e.g. variant calling, assembly). You can trim the first 10-13 bases using **Cutadapt** or **Fastp**.



# Per sequence GC content

## ! Per sequence GC content



A **good** plot of per sequence GC content should be:

- smooth, bell-shaped curve (normal distribution)
- centred roughly around your species' expected GC content (Human 40-41%)
- the peak should match what you expect based on your organism

If your sample has:


- two peaks
- the peak is much higher or lower than you expected
- is very broad and flat.

Something might be wrong with your sample.



# Per sequence GC content

## ! Per sequence GC content



Our sample has two peaks:

1. peak at ~38% and has ~3200 sequences
2. peak at ~53% and has ~6000 sequences


Two peaks usually suggests two different populations of sequences in the data.

Any ideas why???



# Per sequence GC content

## ! Per sequence GC content



The first peak at 38% likely represents our **human brain cells**.

The second peak at 53% might represents:


- contamination from another organism - some bacteria or species have higher GC content
- adapters or technical artifacts - but adapters are usually small and wouldn't cause a whole peak
- mixed library - two different samples got mixed during prep

It is our added ERCC spike-in control...!!! No need to worry.



# Sequence duplication levels

## Sequence Duplication Levels



What a **good** plot looks like?

For genomic DNA sequencing:

- you want most reads to be **unique** (seen only once)
- A small tail of duplicates is acceptable

For RNA-seq:

- **Higher duplication is normal** because highly expressed genes produce many identical reads
- especially true for small transcriptomes like brain-specific libraries

So, no worries for this warning as well!!



# Overrepresented sequences

## ! Overrepresented sequences

Sequence	Count	Percentage	Possible Source
CTTATGTGATAGATGCTCTTTAAATATCTAAGTGCTGGGGTTATGAGT	444	0.37445918479223417	No Hit
CGCTTTGATATTCTCTGCATCCTATTAGGGCTATTGATATTTAACAAAT	396	0.3339771107606413	No Hit
CCGCTTTGATATTCTCTGCATCCTATTAGGGCTATTGATATTTAACAAA	388	0.3272300984220425	No Hit
CGGCTGTCGAGTTGTACGGCCGTTACGCCACGAGTCACGGGGTCTAACGC	382	0.3221698391680934	No Hit
CTTTGATATTCTCTGCATCCTATTAGGGCTATTGATATTTAACAAATAT	381	0.3213264626257685	No Hit
CTGAGACAGAGTCGCTATCGTTATGTCTCTTCCCGGGTCAAGGCGAAA	339	0.28590464784812475	No Hit
GCCTTATGTGATAGATGCTCTTTAAATATCTAAGTGCTGGGGTTATGA	295	0.24879607998583128	No Hit
GTAAGACGCAAGCACCGGCTGTCGAGTTGTACGGCCGTTACGCCACGAGT	259	0.2184345244621366	No Hit
GGAAGCTATATATAGGTGGCTATCTATCCCTACCAAGGCTATATTG	244	0.20578387632726383	No Hit
CTCAGACGCTGCCCTAACTGCGCAGTTAATAATTCTGGCAATTCGTCTCC	227	0.19144647510774138	No Hit
CCCATTTTTAGTTATAATGATGCCCTTATGTGATAGATGCTCTTTAAAT	220	0.1855428393114674	No Hit
GCCTCACTTAGTTACAGTTTTATGGATAATTGGGATATTCTTTGGTATAG	193	0.16277167266869638	No Hit
CCGGGATCGGGCAAGGGGCAACTCGAGCTAATCTCCCGAGCGGCTAGCA	184	0.15518128378777274	No Hit
GTCTCTTCCCGGGTCAAGGCGAAACCGCAGCAAACTTCTCAGACGCT	179	0.15096440107614847	No Hit
GCCATTGCGACCTACCGTAAGCCTATATTTCGTTTTCTGAGACCTATCC	173	0.14590414182219935	No Hit
CGACGAACAACGAAGAGCGACGATGCCGTTTCAGGTGCTCTCAGCGTA	173	0.14590414182219935	No Hit
CCCGTTTCAGGTGGTCTCAGCGTACGGCGGACCTCTGAGAATTGGGAT	160	0.13494024677197627	No Hit
CTCCTTCCCGGGTCAAGGCGAAACCGCAGCAAACTTCTCAGACGCTGC	160	0.13494024677197627	No Hit
GTTATGAGTAGGGATGAGCATAAACCAACTCTCAAGAAGATGGGAA	160	0.13494024677197627	No Hit

It lists any sequences that make up more than 0.1% of your total reads.

It also tries to BLAST them against a small database to guess what they are (e.g., adapters, ribosomal RNA, specific genes).


So, what do you think they could be??

We can also BLAST them to a bigger database like using the NCBI BLAST.



# Adapter content

## ✓ Adapter Content



You should expect there is **no adapter** exists in your data.

Our sample doesn't have adapter contamination which is great.

So, for our sample, we don't have to trim the adapter. But in a lot of scenarios, you need to trim the adapter.

Sometimes, the RNA-seq service provider perform the adapter trimming and filtering for you so you receive a clean data from them.

When the adapter appears, it usually appears at the end of your sequence.





Now we have finished looking at one result from our data, but in total we have 12 FASTQ files we need to look at!!! A lot of work to do, isn't it??

Don't worry, MultiQC can help.

MultiQC is a program that gathers summary statistics from many bioinformatics tools (like FastQC). It combines all the reports from multiple samples into one big, beautiful, interactive HTML report.

Go to directory `~/RNAseq-Workshop/data`.

Run `multiqc .`

Unfortunately, we will encounter an error here because the python version we have here is not compatible to MultiQC.

```
File "/home/vdiuser/miniconda3/envs/RNAseq_test/lib/python3.6/site-packages/ri
ch_click/rich_click.py", line 5, in <module>
    import rich.markdown
File "/home/vdiuser/miniconda3/envs/RNAseq_test/lib/python3.6/site-packages/ri
ch/markdown.py", line 1
    from __future__ import annotations
    ^
SyntaxError: future feature annotations is not defined
```

MultiQC requires python  $\geq 3.8$

But we have python 3.6.10 when we get Conda to solve the environment for us when we install all the software together.

It's okay, let's install a newer version of python and try again.

``conda install -c conda-forge python=3.8``

When it's done, let's try run MultiQC again.

``multiqc .``

```
(RNAseq_env) vdiuser@vdj-33xefq:~/RNAseq-Workshop/data$ multiqc .  
  
/// MultiQC | v1.17  
  
| multiqc | Search path : /home/vdiuser/RNAseq-Workshop/data  
| searching | 100% 36/36  
| fastqc | Found 12 reports  
| multiqc | Report : multiqc_report.html  
| multiqc | Data : multiqc_data  
| multiqc | MultiQC complete  
(RNAseq_env) vdiuser@vdj-33xefq:~/RNAseq-Workshop/data$
```

Seems working... great!!

Then let's use the file explorer to find the “multiqc\_report.html” file and open it.

You can view all 12 results together, and if they are similar, we then can batch processing a standard trimming and cleaning process on them.


## Per Sequence GC Content

12

The average GC content of reads. Normal random library typically have a roughly normal distribution of GC content.

Percentages

Counts



## Sequence Duplication Levels


10

2

The relative level of duplication found for every sequence.

Help

Y-Limits: 0 100



They seem to have similar result which is great!

And if we ever want to trim them or clean them, we can do batch processing.

# Clean up

We don't want our data folder look too messy. So, let's move all the QC result to a new folder.

Let's create a folder “fastqc-result” under “data”:

```
`mkdir fastqc_result`
```

And move all the fastqc result into it:

```
`mv *_fastqc* fastqc_result`
```

```
(RNAseq_env) vdiuser@vdj-33xefq:~/RNAseq-Workshop/data$ mkdir fastqc_result
(RNAseq_env) vdiuser@vdj-33xefq:~/RNAseq-Workshop/data$ mv *_fastqc* fastqc_resu
lt/
(RNAseq_env) vdiuser@vdj-33xefq:~/RNAseq-Workshop/data$ ls
fastqc_result
HBR_Rep1_ERCC-Mix2_Build37-ErccTranscripts-chr22.read1.fastq.gz
HBR_Rep1_ERCC-Mix2_Build37-ErccTranscripts-chr22.read2.fastq.gz
HBR_Rep2_ERCC-Mix2_Build37-ErccTranscripts-chr22.read1.fastq.gz
HBR_Rep2_ERCC-Mix2_Build37-ErccTranscripts-chr22.read2.fastq.gz
HBR_Rep3_ERCC-Mix2_Build37-ErccTranscripts-chr22.read1.fastq.gz
HBR_Rep3_ERCC-Mix2_Build37-ErccTranscripts-chr22.read2.fastq.gz
multiqc_data
multiqc_report.html
UHR_Rep1_ERCC-Mix1_Build37-ErccTranscripts-chr22.read1.fastq.gz
UHR_Rep1_ERCC-Mix1_Build37-ErccTranscripts-chr22.read2.fastq.gz
UHR_Rep2_ERCC-Mix1_Build37-ErccTranscripts-chr22.read1.fastq.gz
UHR_Rep2_ERCC-Mix1_Build37-ErccTranscripts-chr22.read2.fastq.gz
UHR_Rep3_ERCC-Mix1_Build37-ErccTranscripts-chr22.read1.fastq.gz
UHR_Rep3_ERCC-Mix1_Build37-ErccTranscripts-chr22.read2.fastq.gz
(RNAseq_env) vdiuser@vdj-33xefq:~/RNAseq-Workshop/data$
```



# Trimming and filtering FASTQ data

Normally, it is a standard step to trim any adapters or primers that exist in your data and filter out low quality reads.

But our data looks clean so we will skip this step.

Tools for trimming and filtering:

- Cutadapt
- Fastp
- Trimmomatic

That's all for today.. questions??



# Thank you

## Contact us

**Jiajia Li**

Biological Data Science Institute

RN Robertson Building, 46 Sullivan's Creek Rd  
The Australian National University  
Canberra ACT 2600

E [jiajia.li1@anu.edu.au](mailto:jiajia.li1@anu.edu.au)

W <https://bdsi.anu.edu.au/>



Australian  
National  
University

TEQSA PROVIDER ID: PRV12002 (AUSTRALIAN UNIVERSITY)  
CRICOS PROVIDER CODE: 00120C